

```
In [1]: import pandas as pd
import numpy as np
from sklearn import metrics
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: from sklearn.datasets import load_boston
boston = load_boston()
```

```
In [4]: data = pd.DataFrame(boston.data)
```

```
In [5]: data.head()
```

Out[5]:

|   | 0       | 1    | 2    | 3   | 4     | 5     | 6    | 7      | 8   | 9     | 10   | 11     | 12   |
|---|---------|------|------|-----|-------|-------|------|--------|-----|-------|------|--------|------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.98 |
| 1 | 0.02731 | 0.0  | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.14 |
| 2 | 0.02729 | 0.0  | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.03 |
| 3 | 0.03237 | 0.0  | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.94 |
| 4 | 0.06905 | 0.0  | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.33 |

```
In [6]: #Adding the feature names to the dataframe
data.columns = boston.feature_names
data.head()
```

Out[6]:

|   | CRIM    | ZN   | INDUS | CHAS | NOX   | RM    | AGE  | DIS    | RAD | TAX   | PTRATIO | B      | LST/ |
|---|---------|------|-------|------|-------|-------|------|--------|-----|-------|---------|--------|------|
| 0 | 0.00632 | 18.0 | 2.31  | 0.0  | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3    | 396.90 | 4.98 |
| 1 | 0.02731 | 0.0  | 7.07  | 0.0  | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8    | 396.90 | 9.14 |
| 2 | 0.02729 | 0.0  | 7.07  | 0.0  | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8    | 392.83 | 4.03 |
| 3 | 0.03237 | 0.0  | 2.18  | 0.0  | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7    | 394.63 | 2.94 |
| 4 | 0.06905 | 0.0  | 2.18  | 0.0  | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7    | 396.90 | 5.33 |

```
In [7]: #Adding target variable to dataframe
data['PRICE'] = boston.target
```

```
In [8]: #Check the shape of dataframe  
data.shape
```

```
Out[8]: (506, 14)
```

```
In [9]: data.columns
```

```
Out[9]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
              'PTRATIO', 'B', 'LSTAT', 'PRICE'],  
              dtype='object')
```

```
In [10]: data.dtypes
```

```
Out[10]: CRIM      float64  
         ZN        float64  
         INDUS     float64  
         CHAS      float64  
         NOX       float64  
         RM        float64  
         AGE       float64  
         DIS       float64  
         RAD       float64  
         TAX       float64  
         PTRATIO   float64  
         B         float64  
         LSTAT     float64  
         PRICE     float64  
         dtype: object
```

```
In [11]: # Identifying the unique number of values in the dataset  
data.nunique()
```

```
Out[11]: CRIM      504  
         ZN        26  
         INDUS     76  
         CHAS      2  
         NOX       81  
         RM       446  
         AGE      356  
         DIS      412  
         RAD       9  
         TAX      66  
         PTRATIO   46  
         B       357  
         LSTAT    455  
         PRICE    229  
         dtype: int64
```

```
In [12]: # Check for missing values
data.isnull().sum()
```

```
Out[12]: CRIM      0
          ZN        0
          INDUS    0
          CHAS     0
          NOX      0
          RM       0
          AGE      0
          DIS      0
          RAD      0
          TAX      0
          PTRATIO  0
          B        0
          LSTAT    0
          PRICE    0
          dtype: int64
```

```
In [13]: # See rows with missing values
data[data.isnull().any(axis=1)]
```

```
Out[13]:
```

|  | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT | PRICE |
|--|------|----|-------|------|-----|----|-----|-----|-----|-----|---------|---|-------|-------|
|  |      |    |       |      |     |    |     |     |     |     |         |   |       |       |

```
In [14]: # Viewing the data statistics
data.describe()
```

```
Out[14]:
```


|              | CRIM       | ZN         | INDUS      | CHAS       | NOX        | RM         | AGE        |     |
|--------------|------------|------------|------------|------------|------------|------------|------------|-----|
| <b>count</b> | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506 |
| <b>mean</b>  | 3.613524   | 11.363636  | 11.136779  | 0.069170   | 0.554695   | 6.284634   | 68.574901  | 3   |
| <b>std</b>   | 8.601545   | 23.322453  | 6.860353   | 0.253994   | 0.115878   | 0.702617   | 28.148861  | 2   |
| <b>min</b>   | 0.006320   | 0.000000   | 0.460000   | 0.000000   | 0.385000   | 3.561000   | 2.900000   | 1   |
| <b>25%</b>   | 0.082045   | 0.000000   | 5.190000   | 0.000000   | 0.449000   | 5.885500   | 45.025000  | 2   |
| <b>50%</b>   | 0.256510   | 0.000000   | 9.690000   | 0.000000   | 0.538000   | 6.208500   | 77.500000  | 3   |
| <b>75%</b>   | 3.677083   | 12.500000  | 18.100000  | 0.000000   | 0.624000   | 6.623500   | 94.075000  | 5   |
| <b>max</b>   | 88.976200  | 100.000000 | 27.740000  | 1.000000   | 0.871000   | 8.780000   | 100.000000 | 12  |

```
In [15]: # Finding out the correlation between the features
corr = data.corr()
corr.shape
```

```
Out[15]: (14, 14)
```

```
In [16]: # Plotting the heatmap of correlation between features
plt.figure(figsize=(20,20))
sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'
```

Out[16]: <AxesSubplot:>



```
In [17]: # Splitting target variable and independent variables
X = data.drop(['PRICE'], axis = 1)
y = data['PRICE']
```

```
In [18]: # Splitting to training and testing data

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.3, rand
```

```
In [19]: # Import Library for Linear Regression  
from sklearn.linear_model import LinearRegression
```

```
In [20]: # Create a Linear regressor  
lm = LinearRegression()  
  
# Train the model using the training sets  
lm.fit(X_train, y_train)
```

Out[20]: LinearRegression()

```
In [21]: # Value of y intercept  
lm.intercept_
```

Out[21]: 36.357041376594815

```
In [22]: #Converting the coefficient values to a dataframe  
coefficients = pd.DataFrame([X_train.columns, lm.coef_]).T  
coefficients = coefficients.rename(columns={0: 'Attribute', 1: 'Coefficients'})  
coefficients
```

Out[22]:


|    | Attribute | Coefficients |
|----|-----------|--------------|
| 0  | CRIM      | -0.12257     |
| 1  | ZN        | 0.055678     |
| 2  | INDUS     | -0.008834    |
| 3  | CHAS      | 4.693448     |
| 4  | NOX       | -14.435783   |
| 5  | RM        | 3.28008      |
| 6  | AGE       | -0.003448    |
| 7  | DIS       | -1.552144    |
| 8  | RAD       | 0.32625      |
| 9  | TAX       | -0.014067    |
| 10 | PTRATIO   | -0.803275    |
| 11 | B         | 0.009354     |
| 12 | LSTAT     | -0.523478    |

```
In [23]: # Model prediction on train data  
y_pred = lm.predict(X_train)
```


```
In [24]: # Model Evaluation
print('R^2:', metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train, y_pred)) * (len(y_train) - 1) / (len(y_train) - 2))
print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
print('MSE:', metrics.mean_squared_error(y_train, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.7465991966746854
Adjusted R^2: 0.736910342429894
MAE: 3.08986109497113
MSE: 19.07368870346903
RMSE: 4.367343437774162
```


```
In [25]: # Visualizing the differences between actual prices and predicted values
plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```



```
In [26]: # Checking residuals
plt.scatter(y_pred,y_train-y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



```
In [27]: # Checking Normality of errors
sns.distplot(y_train-y_pred)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```



```
In [28]: # Predicting Test data with the model
y_test_pred = lm.predict(X_test)
```

```
In [29]: # Model Evaluation
acc_linreg = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_linreg)
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_test, y_test_pred)) * (len(y_test) - 1) / (len(y_test) - 2))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

R^2: 0.7121818377409185

Adjusted R^2: 0.6850685326005702

MAE: 3.8590055923707407

MSE: 30.05399330712424

RMSE: 5.482152251362985



# Classifying movie reviews: a binary classification example ¶

Two-class classification, or binary classification, may be the most widely applied kind of machine-learning problem. In this example, you'll learn to classify movie reviews as positive or negative, based on the text content of the reviews.

## The IMDB dataset You'll work with the IMDB dataset:

a set of 50,000 highly polarized reviews from the Internet Movie Database. They're split into 25,000 reviews for training and 25,000 reviews for testing, each set consisting of 50% negative and 50% positive reviews.

```
In [1]: from tensorflow.keras.datasets import imdb

# Load the data, keeping only 10,000 of the most frequently occurring words
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=10000)
```

The argument `num_words=10000` means you'll only keep the top 10,000 most frequently occurring words in the training data. Rare words will be discarded. This allows you to work with vector data of manageable size.

```
In [2]: train_data[0]
```

```
Out[2]: [1,
14,
22,
16,
43,
530,
973,
1622,
1385,
65,
458,
4468,
66,
3941,
4,
173,
36,
256,
5,
~
```

```
In [3]: train_labels[0]
```

```
Out[3]: 1
```

Because you're restricting yourself to the top 10,000 most frequent words, no word index will exceed 10,000:

```
In [4]: max([max(sequence) for sequence in train_data])
```

```
Out[4]: 9999
```

```
In [5]: # Let's quickly decode a review

# step 1: Load the dictionary mappings from word to integer index
word_index = imdb.get_word_index()

# step 2: reverse word index to map integer indexes to their respective words
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])

# Step 3: decode the review, mapping integer indices to words
'''Decodes the review. Note that the indices are offset by 3
    because 0, 1, and 2 are reserved indices for "padding," "start of sequence" and "end of sequence" tokens.
    decoded_review = ' '.join([reverse_word_index.get(i-3, '?') for i in train_data[0]])
```

```
In [6]: decoded_review
```

```
Out[6]: "? this film was just brilliant casting location scenery story direction everyone's really suited the part they played and you could just imagine being there robert ? is an amazing actor and now the same being director ? father came from the same scottish island as myself so i loved the fact there was a real connection with this film the witty remarks throughout the film were great it was just brilliant so much that i bought the film as soon as it was released for ? and would recommend it to everyone to watch and the fly fishing was amazing really cried at the end it was so sad and you know what they say if you cry at a film it must have been good and this definitely was also ? to the two little boy's that played the ? of norman and paul they were just brilliant children are often left out of the ? list i think because the stars that play them all grown up are such a big profile for the whole film but these children are amazing and should be praised for what they have done don't you think the whole story was so lovely because it was true and was someone's life after all that was shared with us all"
```

## Preparing the data:

You can't feed lists of integers into a neural network. You have to turn your lists into tensors. There are two ways to do that:

- Pad your lists so that they all have the same length, turn them into an integer tensor of shape (samples, word\_indices), and then use as the first layer in your network a layer

capable of handling such integer tensors (the Embedding layer, which we'll cover in detail later in the book).

- One-hot encode your lists to turn them into vectors of 0s and 1s. This would mean, for instance, turning the sequence [3, 5] into a 10,000-dimensional vector that would be all 0s except for indices 3 and 5, which would be 1s. Then you could use as the first layer in your

```
In [7]: #Encoding the integer sequences into a binary matrix
'''Explanation: I first created 2D matrix of shape(number of examples,10000)
then I looped over each word of each example, if it exist put 1 in its
if not just leave it as 0
ITS JUST ONE HOT ENCODER'''

import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension)) # Creates an all zero m
    for i,sequence in enumerate(sequences):
        results[i,sequence] = 1 # Sets specific indices
    return results

# Vectorize training Data
X_train = vectorize_sequences(train_data)

# Vectorize testing Data
X_test = vectorize_sequences(test_data)
```

```
In [8]: X_train.shape
```

```
Out[8]: (25000, 10000)
```

```
In [9]: #vectorize Labels
y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

```
In [10]: from tensorflow.keras import models
from tensorflow.keras import layers

model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

## Compiling the model

```
In [11]: from tensorflow.keras import optimizers
from tensorflow.keras import losses
from tensorflow.keras import metrics

model.compile(optimizer=optimizers.RMSprop(learning_rate=0.001),
              loss = losses.binary_crossentropy,
              metrics = [metrics.binary_accuracy])
```

## Validating your approach

In order to monitor during training the accuracy of the model on data it has never seen before, you'll create a validation set by setting apart 10,000 samples from the original training data.

```
In [12]: # Input for Validation  
X_val = X_train[:10000]  
partial_X_train = X_train[10000:]  
  
# Labels for validation  
y_val = y_train[:10000]  
partial_y_train = y_train[10000:]
```

```
In [13]: history = model.fit(partial_X_train,  
                             partial_y_train,  
                             epochs=20,  
                             batch_size=512,  
                             validation_data=(X_val, y_val))
```

```
Epoch 1/20
30/30 [=====] - 2s 39ms/step - loss: 0.5154 - binary
_accuracy: 0.7789 - val_loss: 0.3948 - val_binary_accuracy: 0.8479
Epoch 2/20
30/30 [=====] - 0s 13ms/step - loss: 0.3140 - binary
_accuracy: 0.8953 - val_loss: 0.3320 - val_binary_accuracy: 0.8685
Epoch 3/20
30/30 [=====] - 0s 14ms/step - loss: 0.2371 - binary
_accuracy: 0.9179 - val_loss: 0.2880 - val_binary_accuracy: 0.8858
Epoch 4/20
30/30 [=====] - 0s 14ms/step - loss: 0.1897 - binary
_accuracy: 0.9356 - val_loss: 0.2908 - val_binary_accuracy: 0.8832
Epoch 5/20
30/30 [=====] - 0s 13ms/step - loss: 0.1577 - binary
_accuracy: 0.9475 - val_loss: 0.2888 - val_binary_accuracy: 0.8848
Epoch 6/20
30/30 [=====] - 0s 12ms/step - loss: 0.1348 - binary
_accuracy: 0.9557 - val_loss: 0.3059 - val_binary_accuracy: 0.8793
Epoch 7/20
30/30 [=====] - 0s 11ms/step - loss: 0.1131 - binary
_accuracy: 0.9652 - val_loss: 0.3020 - val_binary_accuracy: 0.8847
Epoch 8/20
30/30 [=====] - 0s 12ms/step - loss: 0.0941 - binary
_accuracy: 0.9737 - val_loss: 0.3591 - val_binary_accuracy: 0.8741
Epoch 9/20
30/30 [=====] - 0s 12ms/step - loss: 0.0821 - binary
_accuracy: 0.9773 - val_loss: 0.3465 - val_binary_accuracy: 0.8738
Epoch 10/20
30/30 [=====] - 0s 11ms/step - loss: 0.0698 - binary
_accuracy: 0.9806 - val_loss: 0.3604 - val_binary_accuracy: 0.8749
Epoch 11/20
30/30 [=====] - 0s 12ms/step - loss: 0.0540 - binary
_accuracy: 0.9874 - val_loss: 0.4028 - val_binary_accuracy: 0.8681
Epoch 12/20
30/30 [=====] - 0s 11ms/step - loss: 0.0526 - binary
_accuracy: 0.9858 - val_loss: 0.4011 - val_binary_accuracy: 0.8745
Epoch 13/20
30/30 [=====] - 0s 11ms/step - loss: 0.0391 - binary
_accuracy: 0.9915 - val_loss: 0.4501 - val_binary_accuracy: 0.8752
Epoch 14/20
30/30 [=====] - 0s 11ms/step - loss: 0.0337 - binary
_accuracy: 0.9931 - val_loss: 0.4582 - val_binary_accuracy: 0.8692
Epoch 15/20
30/30 [=====] - 0s 11ms/step - loss: 0.0315 - binary
_accuracy: 0.9921 - val_loss: 0.4697 - val_binary_accuracy: 0.8731
Epoch 16/20
30/30 [=====] - 0s 11ms/step - loss: 0.0282 - binary
_accuracy: 0.9935 - val_loss: 0.4905 - val_binary_accuracy: 0.8741
Epoch 17/20
30/30 [=====] - 0s 11ms/step - loss: 0.0157 - binary
_accuracy: 0.9987 - val_loss: 0.5538 - val_binary_accuracy: 0.8597
Epoch 18/20
30/30 [=====] - 0s 12ms/step - loss: 0.0195 - binary
_accuracy: 0.9961 - val_loss: 0.5443 - val_binary_accuracy: 0.8725
Epoch 19/20
30/30 [=====] - 0s 12ms/step - loss: 0.0145 - binary
_accuracy: 0.9977 - val_loss: 0.5803 - val_binary_accuracy: 0.8626
```

Epoch 20/20

30/30 [=====] - 0s 10ms/step - loss: 0.0086 - binary  
\_accuracy: 0.9997 - val\_loss: 0.7641 - val\_binary\_accuracy: 0.8530

Note that the call to `model.fit()` returns a `History` object. This object has a member `history`, which is a dictionary containing data about everything that happened during training. Let's look at it:


```
In [14]: history_dict = history.history  
history_dict.keys()
```

```
Out[14]: dict_keys(['loss', 'binary_accuracy', 'val_loss', 'val_binary_accuracy'])
```

## Plotting the training and validation loss


```
In [15]: import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [16]: # Plotting Losses  
loss_values = history_dict['loss']  
val_loss_values = history_dict['val_loss']  
  
epochs = range(1, len(loss_values) + 1)  
  
plt.plot(epochs, loss_values, 'bo', label="Training Loss")  
plt.plot(epochs, val_loss_values, 'b', label="Validation Loss")  
  
plt.title('Training and Validation Loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss Value')  
plt.legend()  
  
plt.show()
```



## Plotting the training and validation accuracy

```
In [17]: plt.clf() #Clears the figure
acc_values = history_dict['binary_accuracy']
val_acc_values = history_dict['val_binary_accuracy']
plt.plot(epochs, acc_values, 'bo', label='Training acc')
plt.plot(epochs, val_acc_values, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```





## Retraining a model from scratch

```
In [18]: model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X_train, y_train, epochs=4, batch_size=512)
```

```
Epoch 1/4
49/49 [=====] - 1s 8ms/step - loss: 0.4666 - accuracy: 0.8150
Epoch 2/4
49/49 [=====] - 0s 7ms/step - loss: 0.2716 - accuracy: 0.9035
Epoch 3/4
49/49 [=====] - 0s 7ms/step - loss: 0.2108 - accuracy: 0.9240
Epoch 4/4
49/49 [=====] - 0s 7ms/step - loss: 0.1813 - accuracy: 0.9338
```

```
Out[18]: <keras.callbacks.History at 0x218812dbf10>
```

```
In [19]: results = model.evaluate(X_test, y_test)
```

```
782/782 [=====] - 1s 1ms/step - loss: 0.2846 - accuracy: 0.8868
```

## Using a trained network to generate predictions on new data

```
In [20]: # Making Predictions for testing data
np.set_printoptions(suppress=True)
result = model.predict(X_test)
```

```
782/782 [=====] - 1s 1ms/step
```

```
In [21]: result
```

```
Out[21]: array([[0.20367871],
                [0.9999613 ],
                [0.91047627],
                ...,
                [0.10980374],
                [0.0675016 ],
                [0.67221063]], dtype=float32)
```

```
In [22]: y_pred = np.zeros(len(result))
         for i, score in enumerate(result):
             y_pred[i] = np.asarray([round(x) for x in score])
```

```
In [23]: y_pred
```

```
Out[23]: array([0., 1., 1., ..., 0., 0., 1.])
```

```
In [24]: from tensorflow.keras.metrics import mean_absolute_error
         mae = mean_absolute_error = (y_pred, y_test)
```

```
In [25]: # Error
         mae
```

```
Out[25]: (array([0., 1., 1., ..., 0., 0., 1.]),
          array([0., 1., 1., ..., 0., 0., 0.], dtype=float32))
```

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: %matplotlib inline
```

```
In [3]: fashion_train_df= pd.read_csv('fashion-mnist_train.csv')
```


```
In [4]: fashion_test_df = pd.read_csv('fashion-mnist_test.csv')
```

```
In [5]: fashion_train_df.head()
```

Out[5]:

|   | label | pixel1 | pixel2 | pixel3 | pixel4 | pixel5 | pixel6 | pixel7 | pixel8 | pixel9 | ... | pixel775 | pixel77 |
|---|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|----------|---------|
| 0 | 2     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 0        |         |
| 1 | 9     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 0        |         |
| 2 | 6     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 5      | 0      | ... | 0        |         |
| 3 | 0     | 0      | 0      | 0      | 1      | 2      | 0      | 0      | 0      | 0      | ... | 3        |         |
| 4 | 3     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 0        |         |

5 rows × 785 columns



```
In [6]: fashion_train_df.tail()
```

Out[6]:

|       | label | pixel1 | pixel2 | pixel3 | pixel4 | pixel5 | pixel6 | pixel7 | pixel8 | pixel9 | ... | pixel775 | pi: |
|-------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|----------|-----|
| 59995 | 9     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 0        |     |
| 59996 | 1     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 73       |     |
| 59997 | 8     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 160      |     |
| 59998 | 8     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 0        |     |
| 59999 | 7     | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | 0      | ... | 0        |     |

5 rows × 785 columns



```
In [7]: fashion_train_df.shape
```

Out[7]: (60000, 785)

```
In [8]: fashion_test_df.shape
```

```
Out[8]: (10000, 785)
```

```
In [9]: training = np.array(fashion_train_df,dtype='float32')  
testing = np.array(fashion_test_df,dtype='float32')
```


```
In [10]: training.shape
```

```
Out[10]: (60000, 785)
```

```
In [11]: import random
```

```
In [12]: i = random.randint(0,60001)  
plt.imshow(training[i,1:].reshape(28,28))  
label = training[i,1]  
label
```

```
Out[12]: 0.0
```



```
i = random.randint(0,60001) plt.imshow(training[i,1:].reshape(28,28)) label = training[i,1] label
```

```

In [13]: W_grid = 7
         L_grid = 7

         fig, axes = plt.subplots(L_grid, W_grid, figsize=(17, 17))

         axes = axes.ravel()
         n_training = len(training)

         for i in np.arange(0, W_grid * L_grid):
             index = np.random.randint(0, n_training)
             axes[i].imshow(training[index, 1:].reshape((28, 28)))
             axes[i].set_title(training[index, 0], fontsize=8)
             axes[i].axis('off')

         plt.subplots_adjust(hspace=0.4)

```



```
In [14]: X_train = training[:,1:]/255
y_train = training[:,0]
X_test = testing[:,1:]/255
y_test = testing[:,0]
```

```
In [15]: from sklearn.model_selection import train_test_split
X_train, X_validate, y_train, y_validate = train_test_split(X_train, y_train, t
```

```
In [16]: X_train = X_train.reshape(X_train.shape[0],*(28,28,1))
X_test = X_test.reshape(X_test.shape[0],*(28,28,1))
X_validate = X_validate.reshape(X_validate.shape[0],*(28,28,1))
```

```
In [17]: X_train.shape
```

```
Out[17]: (48000, 28, 28, 1)
```

```
In [18]: X_test.shape
```

```
Out[18]: (10000, 28, 28, 1)
```

```
In [19]: X_validate.shape
```

```
Out[19]: (12000, 28, 28, 1)
```

```
In [20]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D,MaxPooling2D,Dense,Flatten,Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import TensorBoard
```

```
In [21]: cnn_model = Sequential()
cnn_model.add(Conv2D(32,3,3,input_shape = (28,28,1),activation = 'relu'))
cnn_model.add(MaxPooling2D(pool_size= (2,2)))
cnn_model.add(Flatten())
cnn_model.add(Dense(32,activation = 'relu'))
cnn_model.add(Dense(10,activation = 'sigmoid'))
cnn_model.compile(loss = 'sparse_categorical_crossentropy',optimizer = Adam(lea
```

```
In [22]: epochs = 200
```

In [23]:

```
cnn_model.fit(X_train,y_train,batch_size =512,epochs = epochs,verbose = 1,validation_data=(X_test,y_test))
```

```
Epoch 1/200
94/94 [=====] - 2s 12ms/step - loss: 1.3900 - accuracy: 0.5820 - val_loss: 0.7855 - val_accuracy: 0.7256
Epoch 2/200
94/94 [=====] - 1s 9ms/step - loss: 0.6848 - accuracy: 0.7535 - val_loss: 0.6168 - val_accuracy: 0.7744
Epoch 3/200
94/94 [=====] - 1s 9ms/step - loss: 0.5819 - accuracy: 0.7889 - val_loss: 0.5550 - val_accuracy: 0.8021
Epoch 4/200
94/94 [=====] - 1s 9ms/step - loss: 0.5311 - accuracy: 0.8065 - val_loss: 0.5156 - val_accuracy: 0.8148
Epoch 5/200
94/94 [=====] - 1s 9ms/step - loss: 0.5023 - accuracy: 0.8161 - val_loss: 0.4916 - val_accuracy: 0.8238
Epoch 6/200
94/94 [=====] - 1s 8ms/step - loss: 0.4793 - accuracy: 0.8253 - val_loss: 0.4733 - val_accuracy: 0.8345
Epoch 7/200
94/94 [=====] - 1s 8ms/step - loss: 0.4610 - accuracy: 0.8345 - val_loss: 0.4610 - val_accuracy: 0.8345
```

In [24]:

```
evaluation = cnn_model.evaluate(X_test,y_test)
print('Test Accuracy : {:.3f}'.format(evaluation[1]))
```

```
313/313 [=====] - 0s 1ms/step - loss: 0.3384 - accuracy: 0.8835
Test Accuracy : 0.883
```

In [25]:

```
predicted_classes = np.argmax(cnn_model.predict(X_test),axis=-1)
```

```
313/313 [=====] - 0s 947us/step
```

In [26]:

```
predicted_classes
```

```
Out[26]: array([0, 1, 2, ..., 8, 8, 1], dtype=int64)
```

```

In [27]: L = 5
W = 5

fig, axes = plt.subplots(L, W, figsize = (12, 12))
axes = axes.ravel()
for i in np.arange(0, L*W):
    axes[i].imshow(X_test[i].reshape(28, 28))
    axes[i].set_title('Prediction Class:{1} \n true class: {1}'.format(predict
    axes[i].axis('off')
plt.subplots_adjust(wspace = 0.5)

```





```
In [28]: from sklearn.metrics import classification_report

classes = 10
targets = ["Class {}".format(i) for i in range(classes)]
print(classification_report(y_test, predicted_classes, target_names = targets))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Class 0      | 0.86      | 0.81   | 0.83     | 1000    |
| Class 1      | 0.98      | 0.97   | 0.98     | 1000    |
| Class 2      | 0.79      | 0.86   | 0.82     | 1000    |
| Class 3      | 0.88      | 0.89   | 0.89     | 1000    |
| Class 4      | 0.84      | 0.79   | 0.81     | 1000    |
| Class 5      | 0.98      | 0.93   | 0.95     | 1000    |
| Class 6      | 0.69      | 0.71   | 0.70     | 1000    |
| Class 7      | 0.92      | 0.95   | 0.93     | 1000    |
| Class 8      | 0.97      | 0.97   | 0.97     | 1000    |
| Class 9      | 0.94      | 0.96   | 0.95     | 1000    |
| accuracy     |           |        | 0.88     | 10000   |
| macro avg    | 0.88      | 0.88   | 0.88     | 10000   |
| weighted avg | 0.88      | 0.88   | 0.88     | 10000   |

## 1. Problem statement

- We are given Google stock price from 01/2012 to 12/2017.
- The task is to predict the trend of the stock price for 01-06 2018.

## 2. Import library

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dropout
```

## 3. Data processing

### 3.0 import the data

```
In [2]: dataset_train = pd.read_csv('Google_Stock_Price_Train.csv')
```

```
In [3]: dataset_train.head()
```

Out[3]:


|   | Date       | Open   | High   | Low    | Close  | Volume     |
|---|------------|--------|--------|--------|--------|------------|
| 0 | 01/03/2012 | 325.25 | 332.83 | 324.97 | 663.59 | 7,380,500  |
| 1 | 01/04/2012 | 331.27 | 333.87 | 329.08 | 666.45 | 5,749,400  |
| 2 | 01/05/2012 | 329.83 | 330.75 | 326.89 | 657.21 | 6,590,300  |
| 3 | 01/06/2012 | 328.34 | 328.77 | 323.68 | 648.24 | 5,405,900  |
| 4 | 01/09/2012 | 322.04 | 322.29 | 309.46 | 620.76 | 11,688,800 |

```
In [4]: #keras only takes numpy array
training_set = dataset_train.iloc[:, 1: 2].values
```

```
In [5]: training_set.shape
```

Out[5]: (1509, 1)

```
In [6]: plt.figure(figsize=(18, 8))
plt.plot(dataset_train['Open'])
plt.title("Google Stock Open Prices")
plt.xlabel("Time (oldest -> latest)")
plt.ylabel("Stock Open Price")
plt.show()
```



### 3.1 Feature scaling

```
In [7]: import os
if os.path.exists('config.py'):
    print(1)
else:
    print(0)
```

0

```
In [8]: sc = MinMaxScaler(feature_range = (0, 1))
#fit: get min/max of train data
training_set_scaled = sc.fit_transform(training_set)
```

### 3.2 Data structure creation

- taking the reference of past 60 days of data to predict the future stock price.
- It is observed that taking 60 days of past data gives us best results.
- In this data set 60 days of data means 3 months of data.
- Every month as 20 days of Stock price.
- X train will have data of 60 days prior to our date and y train will have data of one day after our date

```
In [9]: ## 60 timesteps and 1 output
X_train = []
y_train = []
for i in range(60, len(training_set_scaled)):
    X_train.append(training_set_scaled[i-60: i, 0])
    y_train.append(training_set_scaled[i, 0])

X_train, y_train = np.array(X_train), np.array(y_train)
```

```
In [10]: X_train.shape
```

```
Out[10]: (1449, 60)
```

```
In [11]: y_train.shape
```

```
Out[11]: (1449,)
```

### 3.3 Data reshaping

```
In [12]: X_train = np.reshape(X_train, newshape =
                               (X_train.shape[0], X_train.shape[1], 1))
```

1. Number of stock prices - 1449
2. Number of time steps - 60
3. Number of Indicator - 1

```
In [13]: X_train.shape
```

```
Out[13]: (1449, 60, 1)
```

## 4. Create & Fit Model

### 4.1 Create model

```
In [14]: regressor = Sequential()
#add 1st Lstm Layer
regressor.add(LSTM(units = 50, return_sequences = True, input_shape = (X_train
regressor.add(Dropout(rate = 0.2))

##add 2nd Lstm Layer: 50 neurons
regressor.add(LSTM(units = 50, return_sequences = True))
regressor.add(Dropout(rate = 0.2))

##add 3rd Lstm Layer
regressor.add(LSTM(units = 50, return_sequences = True))
regressor.add(Dropout(rate = 0.2))

##add 4th Lstm Layer
regressor.add(LSTM(units = 50, return_sequences = False))
regressor.add(Dropout(rate = 0.2))

##add output Layer
regressor.add(Dense(units = 1))
```

```
In [15]: regressor.compile(optimizer = 'adam', loss = 'mean_squared_error')
```

## 4.2 Model fit

```
In [16]: regressor.fit(x = X_train, y = y_train, batch_size = 32, epochs = 100)
```

```
Epoch 1/100
46/46 [=====] - 23s 142ms/step - loss: 0.0276
Epoch 2/100
46/46 [=====] - 7s 146ms/step - loss: 0.0039
Epoch 3/100
46/46 [=====] - 6s 131ms/step - loss: 0.0032
Epoch 4/100
46/46 [=====] - 6s 137ms/step - loss: 0.0032
Epoch 5/100
46/46 [=====] - 7s 143ms/step - loss: 0.0029
Epoch 6/100
46/46 [=====] - 7s 148ms/step - loss: 0.0031
Epoch 7/100
46/46 [=====] - 7s 148ms/step - loss: 0.0028
Epoch 8/100
46/46 [=====] - 7s 146ms/step - loss: 0.0028
Epoch 9/100
46/46 [=====] - 7s 142ms/step - loss: 0.0032
Epoch 10/100
46/46 [=====] - 7s 145ms/step - loss: 0.0027
```

## 4.3 Model evaluation

### 4.3.1 Read and convert

```
In [17]: dataset_test = pd.read_csv('Google_Stock_Price_Test.csv')
```

```
In [18]: dataset_test.head()
```

Out[18]:

|   | Date       | Open        | High        | Low         | Close       | Volume  |
|---|------------|-------------|-------------|-------------|-------------|---------|
| 0 | 02/01/2018 | 1048.339966 | 1066.939941 | 1045.229980 | 1065.000000 | 1237600 |
| 1 | 03/01/2018 | 1064.310059 | 1086.290039 | 1063.209961 | 1082.479980 | 1430200 |
| 2 | 04/01/2018 | 1088.000000 | 1093.569946 | 1084.001953 | 1086.400024 | 1004600 |
| 3 | 05/01/2018 | 1094.000000 | 1104.250000 | 1092.000000 | 1102.229980 | 1279100 |
| 4 | 08/01/2018 | 1102.229980 | 1111.270020 | 1101.619995 | 1106.939941 | 1047600 |

```
In [19]: #keras only takes numpy array
real_stock_price = dataset_test.iloc[:, 1: 2].values
real_stock_price.shape
```

Out[19]: (125, 1)

#### 4.3.2 Concat and convert

```
In [20]: #vertical concat use 0, horizontal uses 1
dataset_total = pd.concat((dataset_train['Open'], dataset_test['Open']),
                           axis = 0)
##use .values to make numpy array
inputs = dataset_total[len(dataset_total) - len(dataset_test) - 60:].values
```

#### 4.3.3 Reshape and scale

```
In [21]: #reshape data to only have 1 col
inputs = inputs.reshape(-1, 1)

#scale input
inputs = sc.transform(inputs)
```

```
In [22]: len(inputs)
```

Out[22]: 185

#### 4.3.4 Create test data strucutre

```
In [23]: X_test = []
         for i in range(60, len(inputs)):
             X_test.append(inputs[i-60:i, 0])
         X_test = np.array(X_test)
         #add dimension of indicator
         X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

```
In [24]: X_test.shape
```

```
Out[24]: (125, 60, 1)
```

### 4.3.5 Model prediction

```
In [25]: predicted_stock_price = regressor.predict(X_test)
```

```
4/4 [=====] - 1s 24ms/step
```

```
In [26]: #inverse the scaled value
         predicted_stock_price = sc.inverse_transform(predicted_stock_price)
```

### 4.3.6 Result visualization

```
In [27]: ##visualize the prediction and real price
         plt.plot(real_stock_price, color = 'red', label = 'Real price')
         plt.plot(predicted_stock_price, color = 'blue', label = 'Predicted price')

         plt.title('Google price prediction')
         plt.xlabel('Time')
         plt.ylabel('Price')
         plt.legend()
         plt.show()
```

